

Yoga Pose Classification



Yogi Bears

Margaret Nesi,
Madison Lease,
Whitney Hannallah



The Problem

01

We are bad at yoga ...



02

Expensive classes and memberships create a barrier for access

- Yoga classes, both online and in-person, studio memberships, and retreats are extremely expensive
- Yoga has many health benefits, if done correctly
- But without instruction, yogis risk injuring themselves

The Process

Phase One

- A. Finding the Data
- B. Creating and training our VGG-16 Model

Up to 93% Accuracy



Phase Two

- A. Working with Tensorflow's MoveNet Thunder
- B. Image Joint Overlay




Phase Three

- A. Display
- B. Classification % Accuracy
- C. Collecting Photos to test User input

Phase One

Trained the VGG-16 CNN classifier
(up to 90% accuracy)

Examined loss function and
tensorboard output



Phase Two




Image display in pop-up

Used MoveNet to find 17 “keypoints” of joints in a human body in an image

Developed a visual to overlay on the image that shows the connections between the 17 keypoints

Hard-coded the title and image path

Phase Three

Calculated percent accuracy of user pose in input image from MoveNet score (displayed above image)

Displayed user image side by side with “perfect pose” both with MoveNet keypoints overlay



Phase Three (cont.)


```
Model: "vgg_base"
```

Layer (type)	Output Shape	Param #
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0

```
=====  
Total params: 14,714,688  
Trainable params: 0  
Non-trainable params: 14,714,688  
Model: "vgg_head"
```

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 25088)	0
dense (Dense)	(None, 256)	6422784
dense_1 (Dense)	(None, 128)	32896
dense_2 (Dense)	(None, 5)	645

```
=====  
Total params: 6,456,325  
Trainable params: 6,456,325  
Non-trainable params: 0
```



Turned run.py (which trains the CNN) into a class

Running movenet.py now both:



A. trains the CNN

B. gives our popup window with overlay

The Final Result & What's Next

What the user will see

Goddess Pose, Percent Accuracy:72.73775916914241



Tree Pose, Percent Accuracy:90.60936577298784



Remaining problem spots

1. Hardcoding the path to the user input image → create a GUI where user can input image
2. Accessing title from run.py
3. Play around with the % accuracy calculation